

Lec18: Greedy Algorithms II

Algorithm I
COMP319-003
Spring 2023

Instructor: Jiyeon Lee

School of Computer Science and Engineering
Kyungpook National University (KNU)

Last time

- Greedy algorithms


- Make a series of choices.
 - Choose this activity, then never backtrack.
- Show that, at each step, your choice does not rule out success.
 - At every step, there exists an optimal solution consistent with the choices we've made so far.
- At the end:
 - You've built only one solution, never having ruled out success,
 - so your solution must be correct.

Today

- Greedy algorithms for **Minimum Spanning Tree**.
- Agenda:
 1. What is a Minimum Spanning Tree?
 2. Short break to introduce some graph theory tools
 3. Prim's algorithm
 4. Kruskal's algorithm

Minimum Spanning Tree

- Say we have an undirected weighted graph.




*A **tree** is a connected graph with no cycles!

- A **spanning tree** is a **tree** that connects all of the vertices.

Minimum Spanning Tree

- Say we have an undirected weighted graph.
- The **cost** of a spanning tree is the sum of the weights on the edges.

This is a
spanning tree.
It has cost 67.




*A **tree** is a connected graph with no cycles!

- A **spanning tree** is a **tree** that connects all of the vertices.

Minimum Spanning Tree

- Say we have an undirected weighted graph.
- The **cost** of a spanning tree is the sum of the weights on the edges.

This is also a
spanning tree.
It has cost 37.




*A **tree** is a connected graph with no cycles!

- A **spanning tree** is a **tree** that connects all of the vertices.

Minimum Spanning Tree

- Say we have an undirected weighted graph.
- The **cost** of a spanning tree is the sum of the weights on the edges.



minimum



of minimal cost




- A **spanning tree** is a **tree** that connects all of the vertices.

Minimum Spanning Tree

- Say we have an undirected weighted graph.
- The **cost** of a spanning tree is the sum of the weights on the edges.

This is a **minimum spanning tree**.
It has cost 37.



minimum



of minimal cost



*A **tree** is a connected graph with no cycles!

- A **spanning tree** is a **tree** that connects all of the vertices.

Why MSTs?

- Network design
 - Connecting cities with roads/electricity/telephone/...
- Cluster analysis
 - e.g., genetic distance
- Image processing
 - e.g., image segmentation
- Useful primitive
 - For other graph algorithms






Figure 2: Fully parsimonious minimal spanning tree of 933 SNPs for 282 isolates of *Y. pestis* colored by location. Morelli et al. Nature genetics 2010

How to find an MST?

- Today, we'll see two greedy algorithms.
- In order to prove that these greedy algorithms work, we'll show something like:
 - *Suppose that our choices so far are consistent with an MST.*
 - *Then the next greedy choice that we make is still consistent with an MST.*


- Again, this is not the only way to prove that these algorithms work.

Let's brainstorm some greedy algorithms!



Brief aside: Cuts in graphs


- A **cut** is a partition of the vertices into two parts:



This is the cut “{A,B,D,E} and {C,I,H,G,F}”

Cuts in graphs


- One or both of the two parts might be disconnected.



This is the cut “{B,C,E,G,H} and {A,D,I,F}”


Cuts in graphs

- This is *not* a cut. Cuts are partitions of vertices.



Let S be a set of edges in G


- We say a cut **respects** S if no edges in S cross the cut.



S is the set of **thick purple** edges

Let S be a set of edges in G


- We say a cut **respects S** if no edges in S cross the cut.
- An edge crossing a cut is called **light** if it has the smallest weight of any edge crossing the cut.



S is the set of **thick purple** edges

Lemma

- Let S be a set of edges, and consider a cut that respects S .
- Suppose there is an MST containing S .
- Let $\{u, v\}$ be a light edge.
- Then there is an MST containing $S \cup \{\{u, v\}\}$




Lemma

- Let S be a set of edges, and consider a cut that respects S .
- Suppose there is an MST containing S .
- Let $\{u, v\}$ be a light edge.
- Then there is an MST containing $S \cup \{\{u, v\}\}$


Aka:

If we haven't ruled out the possibility of success so far, then adding a light edge still won't rule it out.




Proof of Lemma

- Assume that we have:
 - a **cut** that **respects S**




Proof of Lemma

- Assume that we have:
 - a **cut** that respects **S**
 - **S** is part of some **MST T**.
- Say that $\{u, v\}$ is light.
 - lowest cost crossing the cut



Proof of Lemma

- Assume that we have:
 - a **cut** that respects **S**
 - **S** is part of some **MST T**.
- Say that $\{u, v\}$ is light.
 - lowest cost crossing the cut
- If $\{u, v\}$ is in T , we are done.
 - T is an MST containing both $\{u, v\}$ and S .




Proof of Lemma

- Assume that we have:
 - a **cut** that respects **S**
 - **S** is part of some **MST T**.
- Say that $\{u, v\}$ is light.
 - lowest cost crossing the cut
- Say $\{u, v\}$ is not in **T**.
- Note that adding $\{u, v\}$ to **T** will make a cycle.

Claim: Adding any additional edge to a spanning tree will create a cycle.

Proof: Both endpoints are already in the tree and connected to each other.




Proof of Lemma

- Assume that we have:
 - a **cut** that respects **S**
 - **S** is part of some **MST T**.
- Say that $\{u, v\}$ is light.
 - lowest cost crossing the cut
- Say $\{u, v\}$ is not in **T**.
- Note that adding $\{u, v\}$ to **T** will make a cycle.
- There is at least one other edge, $\{x, y\}$, in this cycle crossing the cut.


Claim: Adding any additional edge to a spanning tree will create a cycle.

Proof: Both endpoints are already in the tree and connected to each other.




Proof of Lemma

- Consider swapping $\{u, v\}$ for $\{x, y\}$ in \mathbf{T} .
 - Call the resulting tree \mathbf{T}' .




Proof of Lemma

- Consider swapping $\{u, v\}$ for $\{x, y\}$ in \mathbf{T} .
 - Call the resulting tree \mathbf{T}' .
- **Claim:** \mathbf{T}' is still an MST.
 - It is still a spanning tree (why?)
 - It has cost at most that of \mathbf{T}
 - because $\{u, v\}$ was light.
 - \mathbf{T} had minimal cost.
 - So \mathbf{T}' does too.
- So \mathbf{T}' is an MST containing S and $\{u, v\}$.
 - This is what we wanted.



Lemma

- Let S be a set of edges, and consider a cut that respects S .
- Suppose there is an MST containing S .
- Let $\{u, v\}$ be a light edge.
- Then there is an MST containing $S \cup \{\{u, v\}\}$




Back to MSTs

- How do we find one?
- Today we'll see **two greedy algorithms**.
- The strategy:
 - Make a **series of choices**, adding edges to the tree.
 - Show that each edge we add is **safe to add**:
 - we do not rule out the possibility of success
 - we will choose light edges crossing cuts and use the Lemma.
 - **Keep going** until we have an MST.


Idea 1

- Start growing a tree, greedily add the shortest edge we can to grow the tree.




Idea 1

- Start growing a tree, greedily add the shortest edge we can to grow the tree.




Idea 1

- Start growing a tree, greedily add the shortest edge we can to grow the tree.




Idea 1

- Start growing a tree, greedily add the shortest edge we can to grow the tree.




Idea 1

- Start growing a tree, greedily add the shortest edge we can to grow the tree.




Idea 1

- Start growing a tree, greedily add the shortest edge we can to grow the tree.




Idea 1

- Start growing a tree, greedily add the shortest edge we can to grow the tree.




Idea 1

- Start growing a tree, greedily add the shortest edge we can to grow the tree.



Idea 1

- Start growing a tree, greedily add the shortest edge we can to grow the tree.



We've discovered Prim's algorithm!

- **slowPrim** ($G = (V, E)$, starting vertex s):

- Let (s, u) be the lightest edge coming out of s .
- $MST = \{(s, u)\}$
- $verticesVisited = \{s, u\}$
- **while** $|verticesVisited| < |V|$:
 - find the lightest edge $\{x, v\}$ in E so that:
 - x is in $verticesVisited$
 - v is not in $verticesVisited$
 - add $\{x, v\}$ to MST
 - add v to $verticesVisited$
- **return** MST

At most n iterations
of this while loop.

Time at most m to go
through all the edges
and find the lightest.

Naively, the running time is $O(nm)$:

- For each of $\leq n$ iterations of the while loop:
 - Go through all the edges.

Two questions

1. Does it work?
 - That is, does it actually return a MST?


2. How do we actually implement this?
 - The pseudocode above says “slowPrim”...

Does it work?

- We need to show that our greedy choices **don't rule out success.**
- That is, at every step:
 - If there exists an MST that contains all of the edges S we have added so far...
 - ...then when we make our next choice $\{u, v\}$, there is still an MST containing S and $\{u, v\}$.
- Now it is time to use our lemma!

Lemma


- Let S be a set of edges, and consider a cut that respects S .
- Suppose there is an MST containing S .
- Let $\{u, v\}$ be a light edge.
- Then there is an MST containing $S \cup \{\{u, v\}\}$



Partway through Prim


- Assume that our choices S so far don't rule out success
 - There is an MST consistent with those choices

Q. How can we use our lemma to show that our next choice also does not rule out success?



Partway through Prim

- Assume that our choices **S** so far don't rule out success
 - There is an MST consistent with those choices
- Consider the cut {**visited**, **unvisited**}
 - This cut respects S.




Partway through Prim

- Assume that our choices S so far don't rule out success
 - There is an MST consistent with those choices
- Consider the cut $\{\text{visited}, \text{unvisited}\}$
 - This cut respects S .
- The edge we add next is a **light edge**.
 - Least weight of any edge crossing the cut.
- By the Lemma, **that edge** is safe to add
 - There is still an MST consistent with the new set of edges.

S is the set of edges selected so far.

add this one next



Hooray!

- Our greedy choices **don't rule out success**.
- This is enough (along with an argument by induction) to guarantee correctness of Prim's algorithm.

Formally

- Inductive hypothesis:
 - After adding the t 'th edge, there exists an MST with the edges added so far.
- Base case:
 - After adding the 0'th edge, there exists an MST with the edges added so far. **YEP.**
- Inductive step:
 - If the inductive hypothesis holds for t (aka, the choices so far are safe), then it holds for $t+1$ (aka, the next edge we add is safe).
 - **That's what we just showed.**
- Conclusion:
 - After adding the $n-1$ 'st edge, there exists an MST with the edges added so far.
 - At this point we have a spanning tree, so it better be minimal.


Two questions

1. Does it work?
 - That is, does it actually return a MST?
 - Yes!

2. How do we actually implement this?
 - The pseudocode above says “slowPrim”...


How do we actually implement this?

- Each vertex keeps:
 - the **distance** from itself to the **growing spanning tree**
 - **how to get there.** if you can get there in one edge.




How do we actually implement this?

- Each vertex keeps:
 - the **distance** from itself to the **growing spanning tree**
 - **how to get there.** if you can get there in one edge.
- Choose the closest vertex, add it.




How do we actually implement this?

- Each vertex keeps:
 - the **distance** from itself to the **growing spanning tree**
 - **how to get there.** if you can get there in one edge.
- Choose the closest vertex, add it.



How do we actually implement this?

- Each vertex keeps:
 - the **distance** from itself to the **growing spanning tree**
 - **how to get there.** if you can get there in one edge.
- Choose the closest vertex, add it.
- Update stored info.



Efficient implementation

Every vertex has a key and a parent
Until all the vertices are **reached**:



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) to MST.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) to MST.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) to MST.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) to MST.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) to MST.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) to MST.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) to MST.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) to MST.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) to MST.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) to MST.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) to MST.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) to MST.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) to MST.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

Activate the **unreached** vertex u with the **smallest key**.

for each of u 's unreached neighbors v :

$k[v] = \min(k[v], \text{weight}(u, v))$

if $k[v]$ updated, $p[v] = u$

Mark u as **reached**, and **add**($p[u]$, u) **to MST**.



Can't reach x yet



x is "active"




Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.




This should look pretty familiar

- Very similar to Dijkstra's algorithm!
- **Differences:**
 1. Keep track of $p[v]$ in order to return a tree at the end
 2. Instead of $d[v]$ which we update by
 - $d[v] = \min(d[v], d[u] + w(u, v))$we keep $k[v]$ which we update by
 - $k[v] = \min(k[v], w(u, v))$

Running time

- Exactly the same as Dijkstra:
 - $O(m \log(n))$ using a Red-Black tree as a priority queue.
 - $O(m + n \log(n))$ time if we use a Fibonacci Heap.



Two questions

1. Does it work?
 - That is, does it actually return a MST?
 - Yes!

2. How do we actually implement this?
 - The pseudocode above says “slowPrim”...
 - Implement it basically the same way we’d implement Dijkstra.


What have we learned?

- Prim's algorithm greedily grows a tree
 - similar to Dijkstra's algorithm
- It finds a Minimum Spanning Tree!
 - in time $O(m \log(n))$ if we implement it with a Red-Black Tree.
 - In amortized time $O(m + n \log(n))$ with a Fibonacci heap.
- To prove it worked, we followed the same recipe for greedy algorithms we saw last time.
 - Show that, at every step, we **don't rule out success**.

That's not the only
greedy algorithm for MST!


Other strategy

- what if we just always take **the cheapest edge** whether or not it's connected to what we have so far?




Other strategy

- what if we just always take **the cheapest edge** whether or not it's connected to what we have so far?




Other strategy

- what if we just always take **the cheapest edge** whether or not it's connected to what we have so far?




Other strategy

- what if we just always take **the cheapest edge** whether or not it's connected to what we have so far?




Other strategy

- what if we just always take **the cheapest edge** whether or not it's connected to what we have so far?




Other strategy

- what if we just always take **the cheapest edge** whether or not it's connected to what we have so far?




Other strategy

- what if we just always take the cheapest edge whether or not it's connected to what we have so far?
that won't cause a cycle




Other strategy

- what if we just always take **the cheapest edge** whether or not it's connected to what we have so far?
that won't cause a cycle




Other strategy

- what if we just always take **the cheapest edge** whether or not it's connected to what we have so far?
that won't cause a cycle




Other strategy

- what if we just always take **the cheapest edge** whether or not it's connected to what we have so far?
that won't cause a cycle



Other strategy

- what if we just always take **the cheapest edge** whether or not it's connected to what we have so far?
that won't cause a cycle



We've discovered Kruskal's algorithm!

- **slowKruskal** ($G = (V, E)$):
 - Sort the edges in E by non-decreasing weight.
 - $MST = \{\}$
 - **for** e in E (in sorted order):
 - ← m iterations through this loop
 - **if** adding e to MST won't cause a cycle:
 - ← How do we check this?
 - add e to MST .
 - **return** MST

How **would** you figure out if added e would make a cycle in this algorithm?


Naively, the running time is ???:

- For each of m iterations of the for loop:
 - Check if adding e would cause a cycle...

Two questions


1. Does it work?
 - That is, does it actually return a MST?

2. How do we actually implement this?
 - The pseudocode above says “slowKruskal” ...




At each step, maintaining a forest.

- A **forest** is a collection of disjoint trees




At each step, maintaining a forest.

- A **forest** is a collection of disjoint trees




At each step, maintaining a forest.

- A **forest** is a collection of disjoint trees
- When we add an edge, we **merge two trees**:



At each step, maintaining a forest.


- A **forest** is a collection of disjoint trees
- When we add an edge, we **merge two trees**:



At each step, maintaining a forest.

- A **forest** is a collection of disjoint trees
- When we add an edge, we **merge two trees**:

We never add an edge within a tree since that would create a cycle.




Keep the trees in a special data structure

- **Union-find** data structure also called disjoint-set data structure.
- Used for storing collections of sets
- Supports:
 - **makeSet(u)**: create a set {u}
 - **find(u)**: return the set that u is in
 - **union(u,v)**: merge the set that u is in with the set that v is in.

makeSet(x)
makeSet(y)
makeSet(z)

union(x,y)




Keep the trees in a special data structure

- **Union-find** data structure also called disjoint-set data structure.
- Used for storing collections of sets
- Supports:
 - **makeSet(u)**: create a set {u}
 - **find(u)**: return the set that u is in
 - **union(u,v)**: merge the set that u is in with the set that v is in.

makeSet(x)
makeSet(y)
makeSet(z)

union(x,y)




Keep the trees in a special data structure

- **Union-find** data structure also called disjoint-set data structure.
- Used for storing collections of sets
- Supports:
 - **makeSet(u)**: create a set {u}
 - **find(u)**: return the set that u is in
 - **union(u,v)**: merge the set that u is in with the set that v is in.

makeSet(x)
makeSet(y)
makeSet(z)

union(x,y)

find(x)




Kruskal pseudo-code

- **kruskal**($G = (V, E)$):
 - Sort E by weight in non-decreasing order
 - $MST = \{\}$ /* initialize an empty tree */
 - **for** v in V :
 - **makeSet**(v) /* put each vertex in its own tree in the forest */
 - **for** (u, v) in E : /* go through the edges in sorted order */
 - **if** $\text{find}(u) \neq \text{find}(v)$: /* if u and v are not in the same tree */
 - add (u, v) to MST
 - **union**(u, v) /* merge u 's tree with v 's tree */
 - **return** MST


Once more...

- To start, every vertex is in its own tree.




Once more...

- Then start merging.




Once more...

- Then start merging.




Once more...

- Then start merging.




Once more...

- Then start merging.




Once more...

- Then start merging.




Once more...

- Then start merging.




Once more...

- Then start merging.



Once more...

- Then start merging. **Stop** when we have one big tree!




Running time

- Sorting edges takes $O(m \log(n))$
 - In practice, if the weights are small integers we can use radixSort and take time $O(m)$
- For the rest:
 - n calls to **makeSet** (put each vertex in its own set)
 - $2m$ calls to **find** (for each edge, find its endpoints)
 - $n-1$ calls to **union**
 - (we will never add more than $n-1$ edges to the tree, so we will never call **union** more than $n-1$ times.)
- Total running time:
 - Worst-case $O(m \log(n))$, just like Prim with a RBtree.
 - Closer to $O(m)$ if you can do radixSort

Two questions

1. Does it work?

- That is, does it actually return a MST?



2. How do we actually implement this?


- The pseudocode above says “slowKruskal”...
 - Worst-case running time $O(m \log(n))$ using a union-find data structure.

Does it work?

- We need to show that our greedy choices **don't rule out success**.
- That is, at every step:
 - There exists an MST that contains all of the edges we have added so far.
- Now it is time to use our lemma again!

Lemma


- Let S be a set of edges, and consider a cut that respects S .
- Suppose there is an MST containing S .
- Let $\{u, v\}$ be a light edge.
- Then there is an MST containing $S \cup \{\{u, v\}\}$



Partway through Kruskal


- Assume that our choices **S** so far don't rule out success.
 - There is an MST extending them
- The **next edge** we add will merge two trees, **T1, T2**

S is the set of
edges selected so far.




Partway through Kruskal

- Assume that our choices **S** so far don't rule out success.
 - There is an MST extending them
- The **next edge** we add will merge two trees, **T1**, **T2**



Partway through Kruskal


- Assume that our choices **S** so far don't rule out success.
 - There is an MST extending them
- The **next edge** we add will merge two trees, **T1**, **T2**
- Consider the cut $\{T1, V - T1\}$.
 - This cut respects **S**
 - Our **new edge is light** for the cut



Partway through Kruskal

- Assume that our choices **S** so far don't rule out success.
 - There is an MST extending them
- The **next edge** we add will merge two trees, **T1**, **T2**
- Consider the cut $\{\mathbf{T1}, \mathbf{V - T1}\}$.
 - This cut respects **S**
 - Our **new edge is light** for the cut
- By the Lemma, **that edge** is safe to add.
 - There is still an MST extending the new set

S is the set of edges selected so far.



Hooray!

- Our greedy choices **don't rule out success**.
- This is enough (along with an argument by induction) to guarantee correctness of Kruskal's algorithm.

Formally

- Inductive hypothesis:
 - After adding the t 'th edge, there exists an MST with the edges added so far.
- Base case:
 - After adding the 0'th edge, there exists an MST with the edges added so far. **YEP.**
- Inductive step:
 - If the inductive hypothesis holds for t (aka, the choices so far are safe), then it holds for $t+1$ (aka, the next edge we add is safe).
 - **That's what we just showed.**
- Conclusion:
 - After adding the $n-1$ 'st edge, there exists an MST with the edges added so far.
 - At this point we have a spanning tree, so it better be minimal.

Two questions

1. Does it work?
 - That is, does it actually return a MST?
 - Yes!
2. How do we actually implement this?
 - The pseudocode above says “slowKruskal”...
 - Using a union-find data structure!

What have we learned?

- Kruskal's algorithm greedily grows a forest
- It finds a Minimum Spanning Tree in time $O(m \log(n))$
 - if we implement it with a Union-Find data structure
 - if the edge weights are reasonably-sized integers and we ignore the inverse Ackerman function, basically $O(m)$ in practice.
- To prove it worked, we followed the same recipe for greedy algorithms we saw last time.
 - Show that, at every step, we **don't rule out success**.

Compare and contrast

- **Prim:**

- Grows a tree.
- Time $O(m \log(n))$ with a red-black tree
- Time $O(m + n \log(n))$ with a Fibonacci heap

Prim might be a better idea
on dense graphs if you can't
radixSort edge weights


- **Kruskal:**

- Grows a forest.
- Time $O(m \log(n))$ with a union-find data structure
- If you can do radixSort on the edge weights, morally $O(m)$

Kruskal might be a better
idea on sparse graphs if you
can radixSort edge weights

Both Prim and Kruskal

- Greedy algorithms for MST.
- Similar reasoning:
 - Optimal substructure: subgraphs generated by cuts.
 - The way to make safe choices is to choose light edges crossing the cut.



Recap

- Two algorithms for Minimum Spanning Tree
 - Prim's algorithm
 - Kruskal's algorithm
- Both are examples of **greedy algorithms!**
 - Make a series of choices.
 - Show that at each step, your choice does not rule out success.
 - At the end of the day, you haven't ruled out success, so you must be successful.

An aerial night photograph of a city. In the center, a large, modern building with a prominent dome and many windows is illuminated. To the left, a tall, slender water tower stands out against the dark sky. In the foreground, a winding road with light trails from cars leads towards the building. A lake with a small fountain is visible in the lower-left corner. The background shows a dense urban area with various buildings and distant mountains under a dark sky. The text "Any Question?" is overlaid in the center in a large, white, bold font.

Any Question?